Krishna Saketh K

Data Engineer/Data Analyst
+1 850-980-5076 | krishnasakethai@gmail.com
https://www.linkedin.com/in/ krishna-saketh-kamadana | https://github.com/ Kr1shna304

Professional Summary

- 6+ years of experience in Data Engineering, Analytics Engineering, Business Intelligence, and Cloud Data Platforms across government, higher education, consulting, retail, financial services, and banking domains.
- Strong expertise in designing and developing scalable ETL/ELT pipelines using Python, SQL, PySpark, Azure Data Factory, Databricks, Snowflake, and dbt for large-scale structured and semi-structured data processing.
- Hands-on experience building modern data lake, lakehouse, and data warehouse solutions leveraging Azure, AWS, Delta Lake, Snowflake, and distributed processing frameworks.
- Proven track record of optimizing high-volume data pipelines through partitioning strategies, performance tuning, data quality frameworks, and workflow orchestration using Airflow and cloud-native services.
- Experience delivering enterprise data migration and modernization initiatives, including Oracle-to-Azure migrations, data reconciliation, governance, lineage, and regulatory compliance requirements.
- Skilled in developing real-time and batch data processing solutions using Apache Spark, Kafka, event-driven architectures, and cloud-based data integration platforms.
- Experience supporting advanced analytics and AI initiatives through Retrieval-Augmented Generation (RAG), vector search pipelines, semantic search, and preparation of feature-ready datasets for machine learning workloads.
- Strong collaboration experience working with business stakeholders, analysts, data scientists, and leadership teams to deliver reliable data platforms, reporting solutions, and data-driven decision-making capabilities.

Technical Skills

- **Programming & Query Languages:** Python, PySpark, SQL, Spark SQL, PL/SQL, T-SQL, Shell Scripting
- **Data Engineering, ETL & Streaming:** Apache Spark, Apache Kafka, Spark Streaming, Apache Airflow, Alteryx, Informatica PowerCenter, SSIS, dbt, CDC (Change Data Capture), ETL/ELT Pipelines, Data Ingestion Frameworks, Event-Driven Architecture, REST APIs, Workflow Orchestration, Incremental Loading, Batch Processing, Data Migration, Schema Evolution, Schema Mapping, Data Transformation, Data Validation, Data Reconciliation, Production Support
- **Cloud & Data Platforms:** Azure Data Factory (ADF), Azure Databricks, ADLS Gen2, Azure Synapse Analytics, Azure SQL Database, Azure OpenAI Service, Azure DevOps, AWS S3, EMR, Glue, Lambda, CloudWatch, IAM, Google BigQuery, Cloud Storage, Snowflake, Delta Lake
- **Data Warehousing, Modeling & Governance:** Dimensional Modeling, Star Schema, Snowflake Schema, Data Warehouse Design, Data Lake Architecture, Lakehouse Architecture, Data Lineage, Metadata Management, Data Governance, Unity Catalog, Microsoft Purview, AWS Glue Catalog, Data Cataloging, Audit & Compliance Reporting
- **Databases & Storage:** Oracle, SQL Server, PostgreSQL, MySQL, MongoDB
- **Big Data & File Formats:** Hadoop, HDFS, Hive, YARN, Distributed Computing, Parquet, Avro, JSON, CSV, XML
- **DevOps & Development Tools:** Git, GitHub, Jenkins, Docker, Kubernetes, Azure DevOps Pipelines, CI/CD Pipelines, Linux/Unix, JIRA
- **Generative AI & Vector Search:** OpenAI APIs, LangChain, Retrieval-Augmented Generation (RAG), Prompt Engineering, Embedding Models, Vector Embeddings, Semantic Search, Vector Search, Document Retrieval, Knowledge Retrieval Systems, FAISS, Pinecone
- **BI & Analytics:** Tableau, Power BI, Business Intelligence, Dashboard Development, KPI Reporting, Self-Service Analytics, Reporting & Analytics, Operational Reporting

Work History

Florida Commerce, Tallahassee, FL **Mar 2026 – Present**
Data Engineer II

- Designed and maintained ETL/ELT pipelines using Alteryx, Python, and Azure Data Factory to integrate workforce, grant, and economic datasets from 15+ state agency systems, supporting centralized reporting and analytics initiatives.
- Developed PySpark transformations on Azure Databricks to process 50M+ records monthly, optimizing partitioning and file management strategies that reduced pipeline execution time by 35%.
- Built and maintained Snowflake reporting datasets and dbt models, implementing incremental transformations and automated data quality checks across 100+ analytics tables, reducing data validation effort by 40%.

Florida Highway Safety and Motor Vehicles, Tallahassee, FL **Oct 2025 – Feb 2026**
Data Engineer

- Contributed to the migration of legacy Oracle-based licensing, vehicle registration, and enforcement datasets to an Azure-based analytics platform (ADF, ADLS Gen2, Synapse), designing extraction, transformation, and schema-mapping processes while maintaining compatibility with downstream reporting and operational systems.
- Developed Azure Data Factory pipelines to extract and load data from Oracle systems, processing 20M+ records during migration and validation cycles.

Florida State University, Tallahassee, FL **Nov 2023 – Sep 2025**
Data Engineer

- Designed and maintained enterprise-scale data pipelines on Azure-centric data platform (ADF, ADLS Gen2, Synapse) integrating academic, administrative, and research datasets into a centralized analytics system supporting multiple university departments.
- Built and optimized PySpark-based ETL workflows on Databricks, handling large structured and semi-structured datasets, and resolving recurring performance issues caused by data skew and uneven partition distribution in batch processing jobs.
- Improved pipeline reliability by redesigning ingestion workflows in Azure Data Factory, introducing incremental load strategies and dependency-based orchestration to address frequent failures caused by late-arriving and inconsistent source data.
- Developed cloud data lake architecture on Azure ADLS Gen2, defining standardized folder structures, partitioning strategy, and lifecycle policies to reduce query scan costs and improve downstream analytics performance.

Environment: Azure Data Factory (ADF), Azure Data Lake Storage Gen2 (ADLS Gen2), Azure Synapse Analytics, Azure Databricks, PySpark, Apache Spark, Python, SQL, Delta Lake, Apache Kafka, REST APIs, JSON, Parquet, Azure DevOps, Git, Linux/Unix, ETL/ELT Pipelines, Data Lake ArchitectureOpenAI API, LangChain, FAISS, Vector Search.

Tiger Analytics, India **May 2021 – May 2023**
Big Data Engineer

- Designed and optimized Spark-based batch processing pipelines for large-scale retail and financial datasets (~20M+ records/day), addressing recurring data skew issues in join-heavy transformations by implementing partition tuning and
- Worked across retail and financial services client verticals simultaneously, applying domain-specific data modeling patterns — transaction fact tables, customer dimension hierarchies — to deliver analytics-ready datasets meeting client SLA requirements.
- Documented end-to-end data flow architecture, transformation logic, and dependency maps for all major pipelines, reducing onboarding time for new engineers and eliminating single points of knowledge failure on the team.
- Collaborated with infrastructure teams to right-size EMR cluster configurations — instance types, executor memory, and core counts — reducing average cluster compute cost by approximately 20% without degrading job SLAs.
- Created pipeline execution monitoring dashboards in CloudWatch, tracking job run durations, failure rates, and data volume trends, giving engineering and stakeholder teams real-time visibility into pipeline health.

Environment: Apache Spark, PySpark, Spark SQL, Python, SQL, Hadoop (HDFS, Hive, YARN), AWS (S3, EMR, Glue), Apache Kafka, REST APIs, JSON, Parquet, Avro, ETL/ELT Pipelines, Data Lake Architecture, Distributed Computing, Git, Jenkins, Linux/Unix, Tableau, Power BI, Performance Tuning (Partitioning, Broadcast Joins, Caching), Agile (Scrum)

LatentView Analytics, India **Jan 2019 – May 2021**
Data Analyst/ETL Developer

- Developed and maintained SQL-based data extraction and transformation pipelines for banking transactions and customer account datasets, processing over 5M+ daily records using SQL Server and Oracle databases.
- Built and optimized ETL workflows using SSIS and Informatica PowerCenter, improving data processing efficiency by ~30% through query tuning and reduced redundant transformations.
- Performed data validation, reconciliation, and quality checks across core banking datasets including transactions, loans, and credit card systems to ensure data accuracy and regulatory compliance.

Education

- Master of Science (MS) in Computer Science – Florida State University, Tallahassee, FL, USA